

# DBMS Architecture

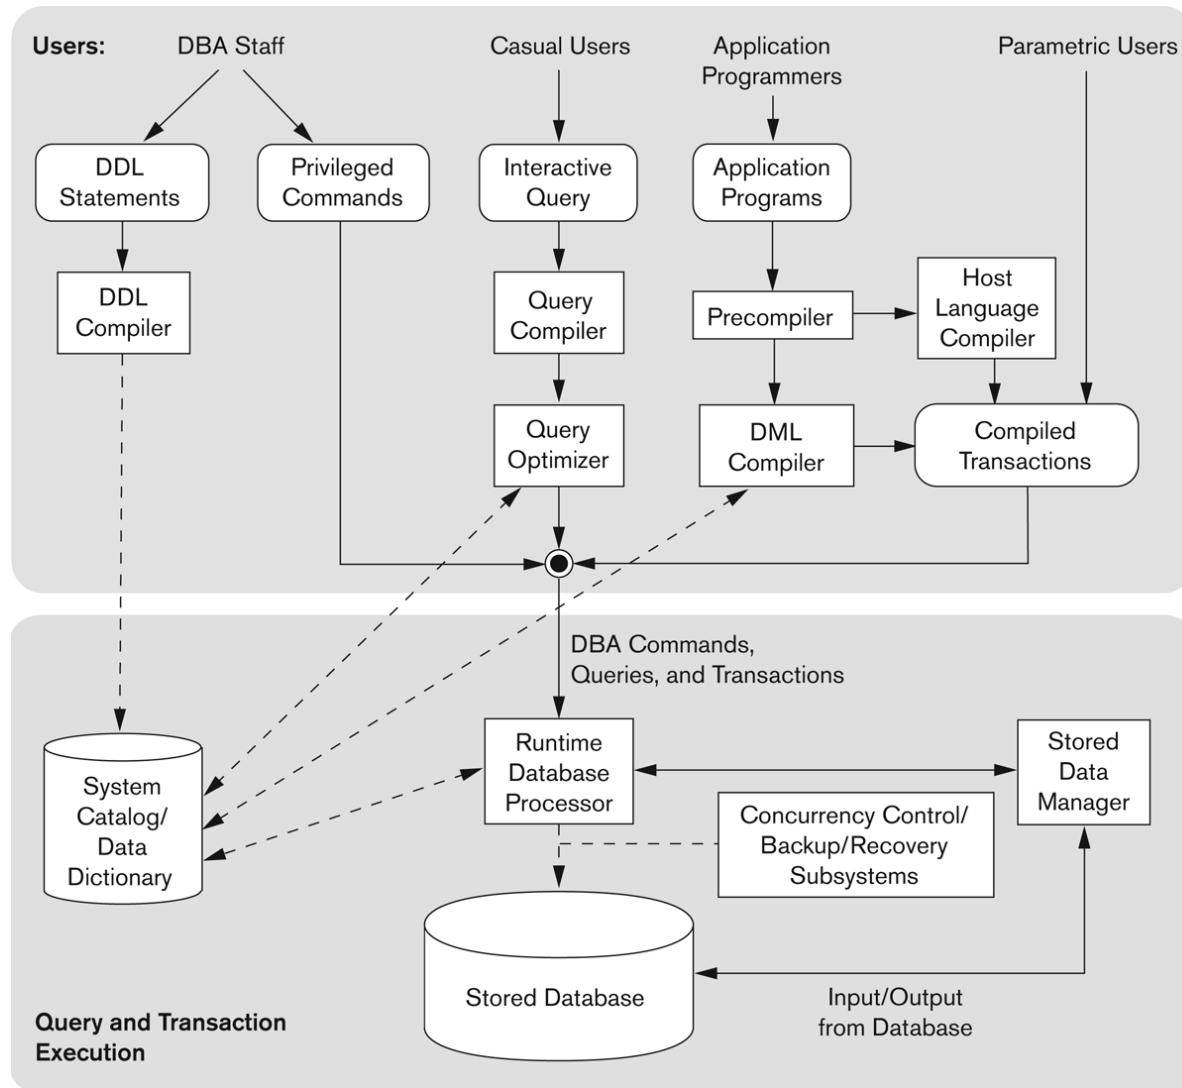
Prof. Daniel S. Kaster

Departamento de Computação  
Universidade Estadual de Londrina  
dskaster@uel.br

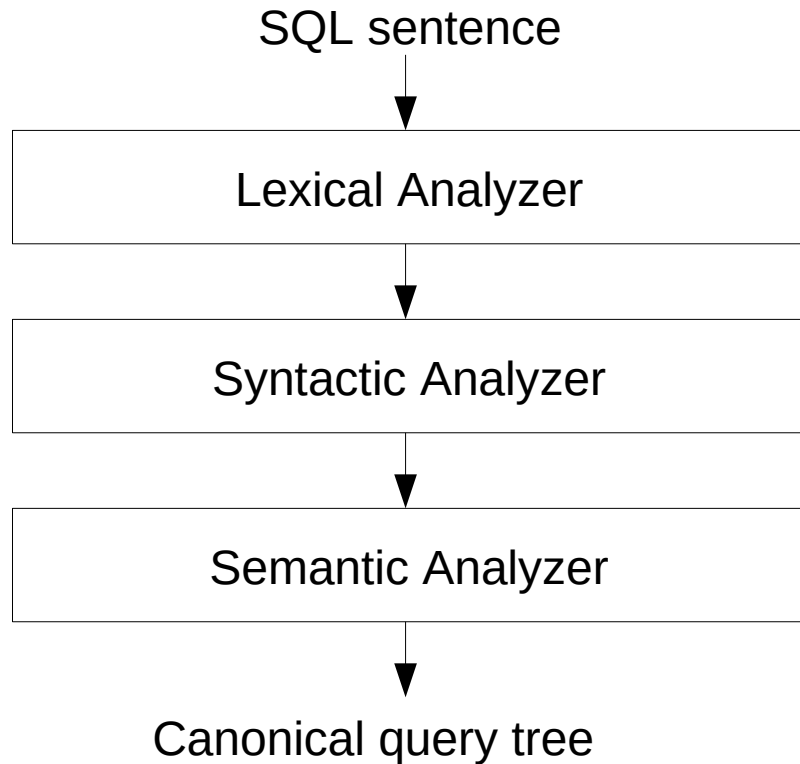
# Outline

- Typical DBMS modules
- Query compilation and optimization
- Storage and indexing
- Query operator algorithms
  - **Select**
- Transaction processing and concurrency control
- Failure recovery

# Typical DBMS Component Modules



# Compilador



Query parsing (tokens)

Language syntax analysis

Verification of objects and identifiers (tables, attributes, etc.), type check, parameters, etc.



Accesses the DBMS catalog

# Query Tree

- A query tree is a representation of an SQL sentence as a tree of operators, in which:
  - Leaves are the input relations;
  - Internal nodes are the operators;
  - The execution occurs bottom-up, where each operation generates an intermediate result (table), which serves as input to the operation directly above in the query tree.
- The canonical query tree directly corresponds to how the query was written by the user

# Example

- Consider the following query over the company example:

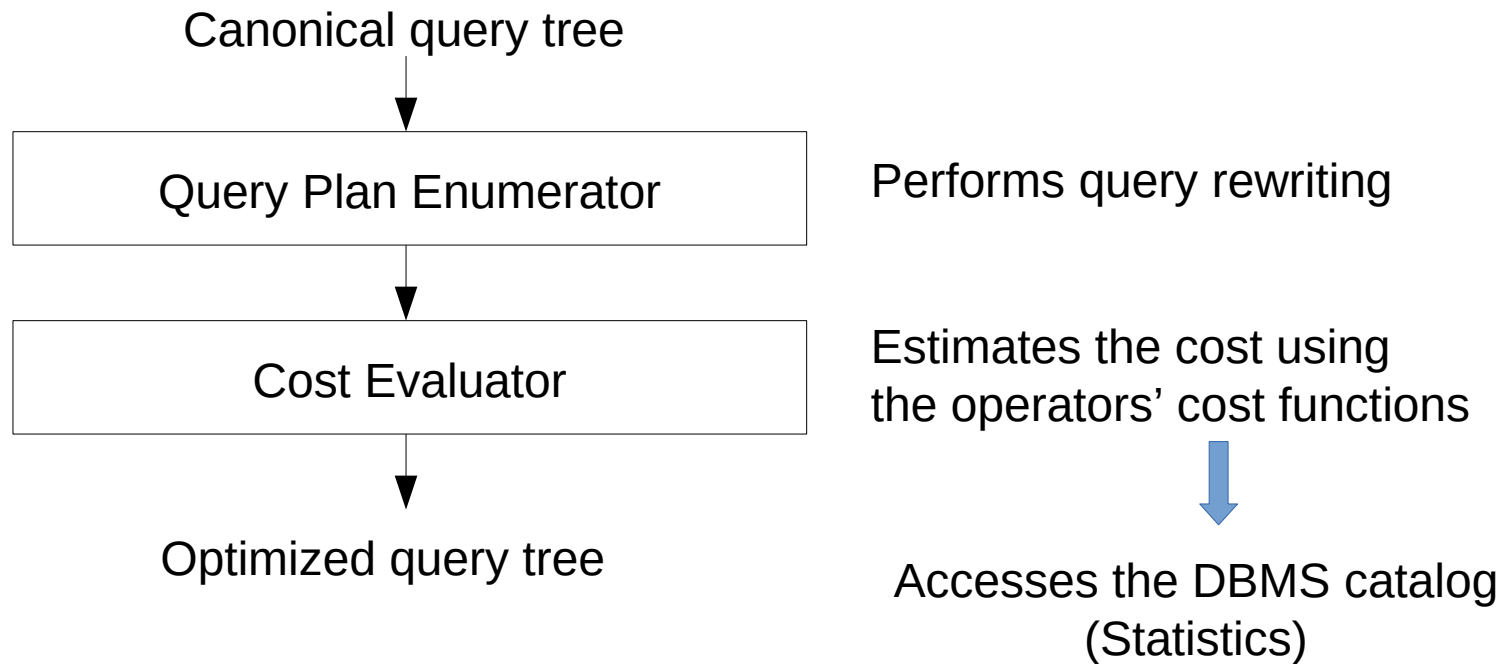
```
SELECT e.name  
FROM employee e, works_on w, project p  
WHERE p.pname = 'aquarius' AND p.pnum =  
w.pnum  
AND e.ssn = w.ssn AND e.bdate >= '1957-12-31';
```

- Notice: pnum, pname and ssn are unique

# Query Optimization

- The “optimization” step aims to generate a query tree equivalent to the initial one, but with more efficient execution
- Didactically, it consists of two parts:
  - Logical optimization: aims to explore the algebraic properties of relational queries
  - Physical optimization: aims to identify which algorithms can perform each operation most efficiently
    - Although the optimization process performs logical and physical optimization together

# The Query Optimizator





# Query Plan Enumerator

- The optimization step has to generate an execution tree equivalent to the initial one
  - To ensure producing the same result
- Therefore, this processing is limited to the transformations possible in a relational query
- According to the meaning of each operation, a set of transformation rules was defined
  - Following the strong mathematical foundation of the relational model

# Equivalence Rules

- R1) Cascading  $\sigma$

$$\sigma_{c1 \text{ AND } c2 \text{ AND } \dots \text{ AND } c_n}(R) \equiv \sigma_{c1}(\sigma_{c2}(\dots(\sigma_{c_n}(R))\dots))$$

- R2) Commutativity of  $\sigma$

$$\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$$

- R3) Cascading  $\pi$

$$\pi_{\text{list1}}(\pi_{\text{list2}}(\dots(\pi_{\text{listn}}(R))\dots)) \equiv \pi_{\text{list1}}(R)$$

- R4) Commutativity between  $\sigma$  and  $\pi$

$$\pi_{A1, A2, \dots, A_n}(\sigma_c(R)) \equiv \sigma_c(\pi_{A1, A2, \dots, A_n}(R))$$

- R5) Commutativity of  $\times$  and  $\bowtie$

$$R \times S \equiv S \times R \quad \text{e} \quad R \bowtie_c S \equiv S \bowtie_c R$$

# Equivalence Rules(2)

- R6) Commutativity of  $\sigma$  with  $\bowtie$  (or  $\times$ ,  $\cup$ ,  $\cap$ )

$$\sigma_c(R \bowtie S) \equiv \sigma_c(R) \bowtie S \text{ ou } \sigma_c(R \bowtie S) \equiv (\sigma_{c1}(R)) \bowtie (\sigma_{c2}(S))$$

- R7) Commutativity of  $\pi$  with  $\bowtie$  (or  $\times$ )

$$\pi_L(R \bowtie S) \equiv \pi_L((\pi_{A1, \dots, An, A1', \dots, An'}(R)) \bowtie (\pi_{B1, \dots, Bn, B1', \dots, Bn'}(S)))$$

- R8) Commutativity of set operators

$\cup$  and  $\cap$  are commutative, but  $-$  is not

- R9) Associativity of  $\cup$ ,  $\cap$ ,  $\times$  and  $\bowtie$

Let  $\theta$  be the operator:  $(R \theta S) \theta T \equiv R \theta (S \theta T)$

- R10) Conversion of  $(\sigma_c(R \times S))$  into  $(R \bowtie_c S)$

# Logical Optimization

- Logical optimization performs what is called query rewriting, aiming to obtain a more efficient query tree
- Fundamentals of logical optimization algorithms:
  - Apply SELECT and PROJECT operations whenever possible before binary operations, as binary operations are generally multiplicative functions on input sizes;
  - Replace sets of equivalent operations with less costly operations;
  - Define a suitable order for executing binary operations, as the order can significantly change the cost.

# A Generic Optimization Algorithm

- 1) Break conjunctive selections into a cascade of selections, to allow moving the selections (R1)
- 2) Push down selections whenever useful (R2, R4, R6)
- 3) Convert  $\sigma_c(R \times S)$  into  $(R \bowtie S)$  (R10)
- 4) Rearrange leaf nodes to optimize binary operations, positioning relations with more restrictive selection operations as leaves (R5, R8, R9)
- 5) Push down projections in the query tree (R3, R4, R7)

# Physical Optimization

- The physical optimization step aims to identify the most efficient algorithms that can be used to execute each operator in the query tree to generate a physical execution plan with minimum cost
- Each algorithm has requirements to be selected
  - For example, the existence of indexes, the organization of the data file, the amount of memory available for the algorithm, etc.
- The positioning of an operator in the query tree during optimization impacts the availability of algorithms for the operator, tying together logical and physical optimizations
  - For example, you can only use indices of operands that are leaves of the execution tree (i.e., original relations)

# Algorithms for Query Operators

- Many query operators have several alternative algorithms to execute the operator
- Physical optimization involves selecting the cheapest algorithm for each case
  - The cost of the algorithms is usually stated based on the number of disk accesses (or number of blocks accessed)

# Algorithms: Select

- A trivial strategy to implement the select operator performs a sequential scan on the input relation and adds the tuples satisfying the condition to the output relation

```
seqScanSelect(Relation R, Condition condition) {  
    for each tuple r in R do {  
        if (condition(r)) {  
            result.add(r);  
        }  
    }  
    return result;  
}
```

- However, this strategy is not suitable for accessing the secondary memory

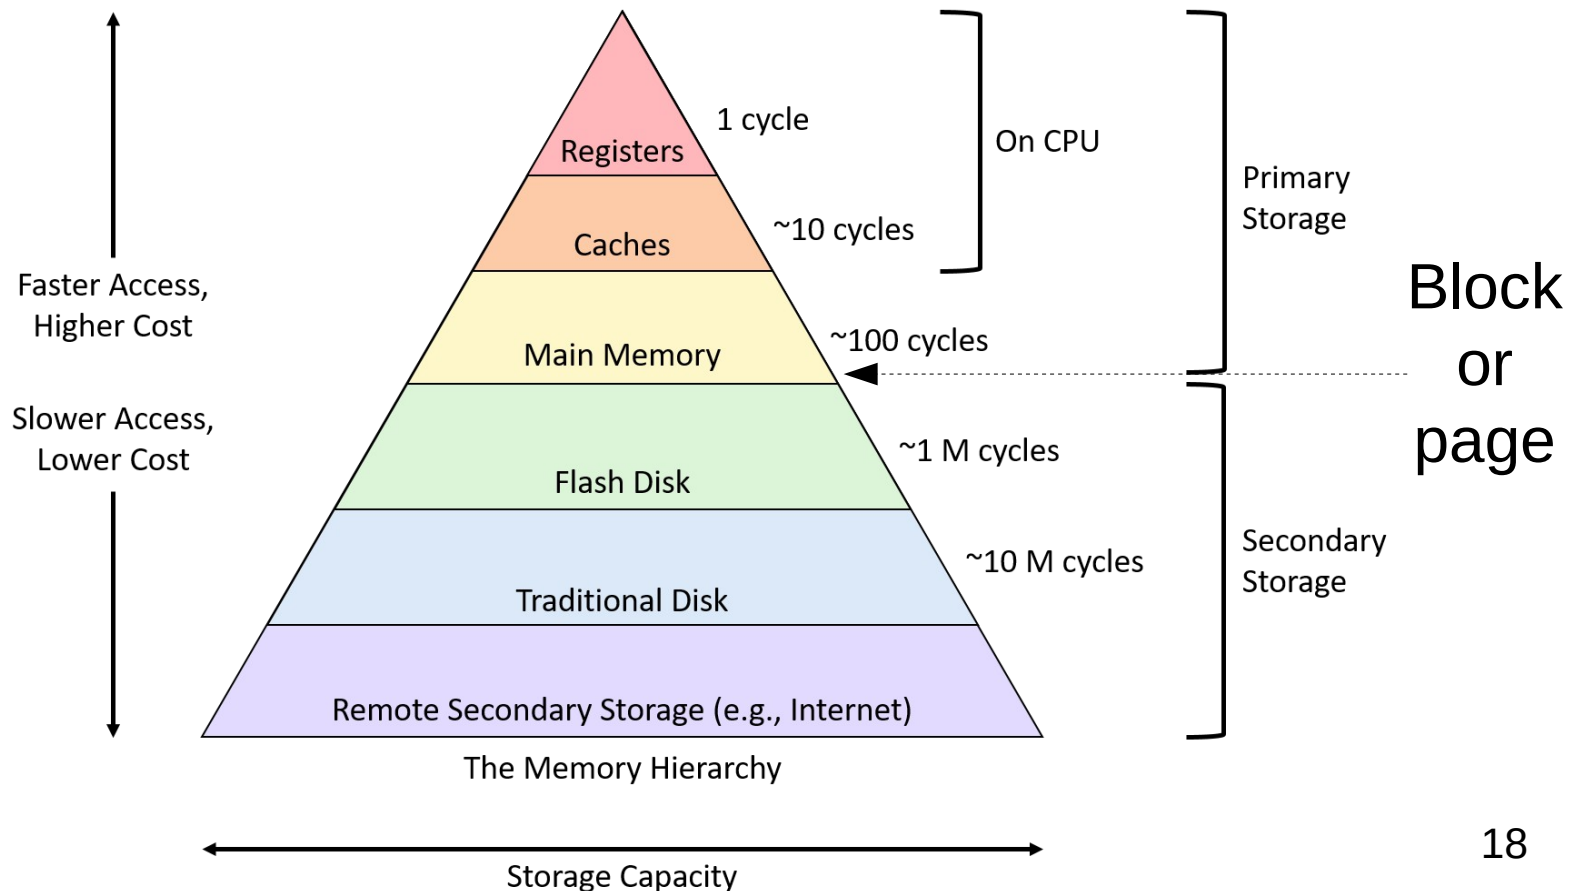


# DBMS Storage

- A database is stored as a collection of files
  - Each file is a sequence of records
  - A record is a sequence of values for fields/attributes
- This material considers the following simplification
  - Each persistent database object (table, index, etc.) is stored in a different file
  - Each file has records of the same data type

# Data Transfer in the Memory Hierarchy

- Memory hierarchy relies on transfer units between levels to ease block replacement



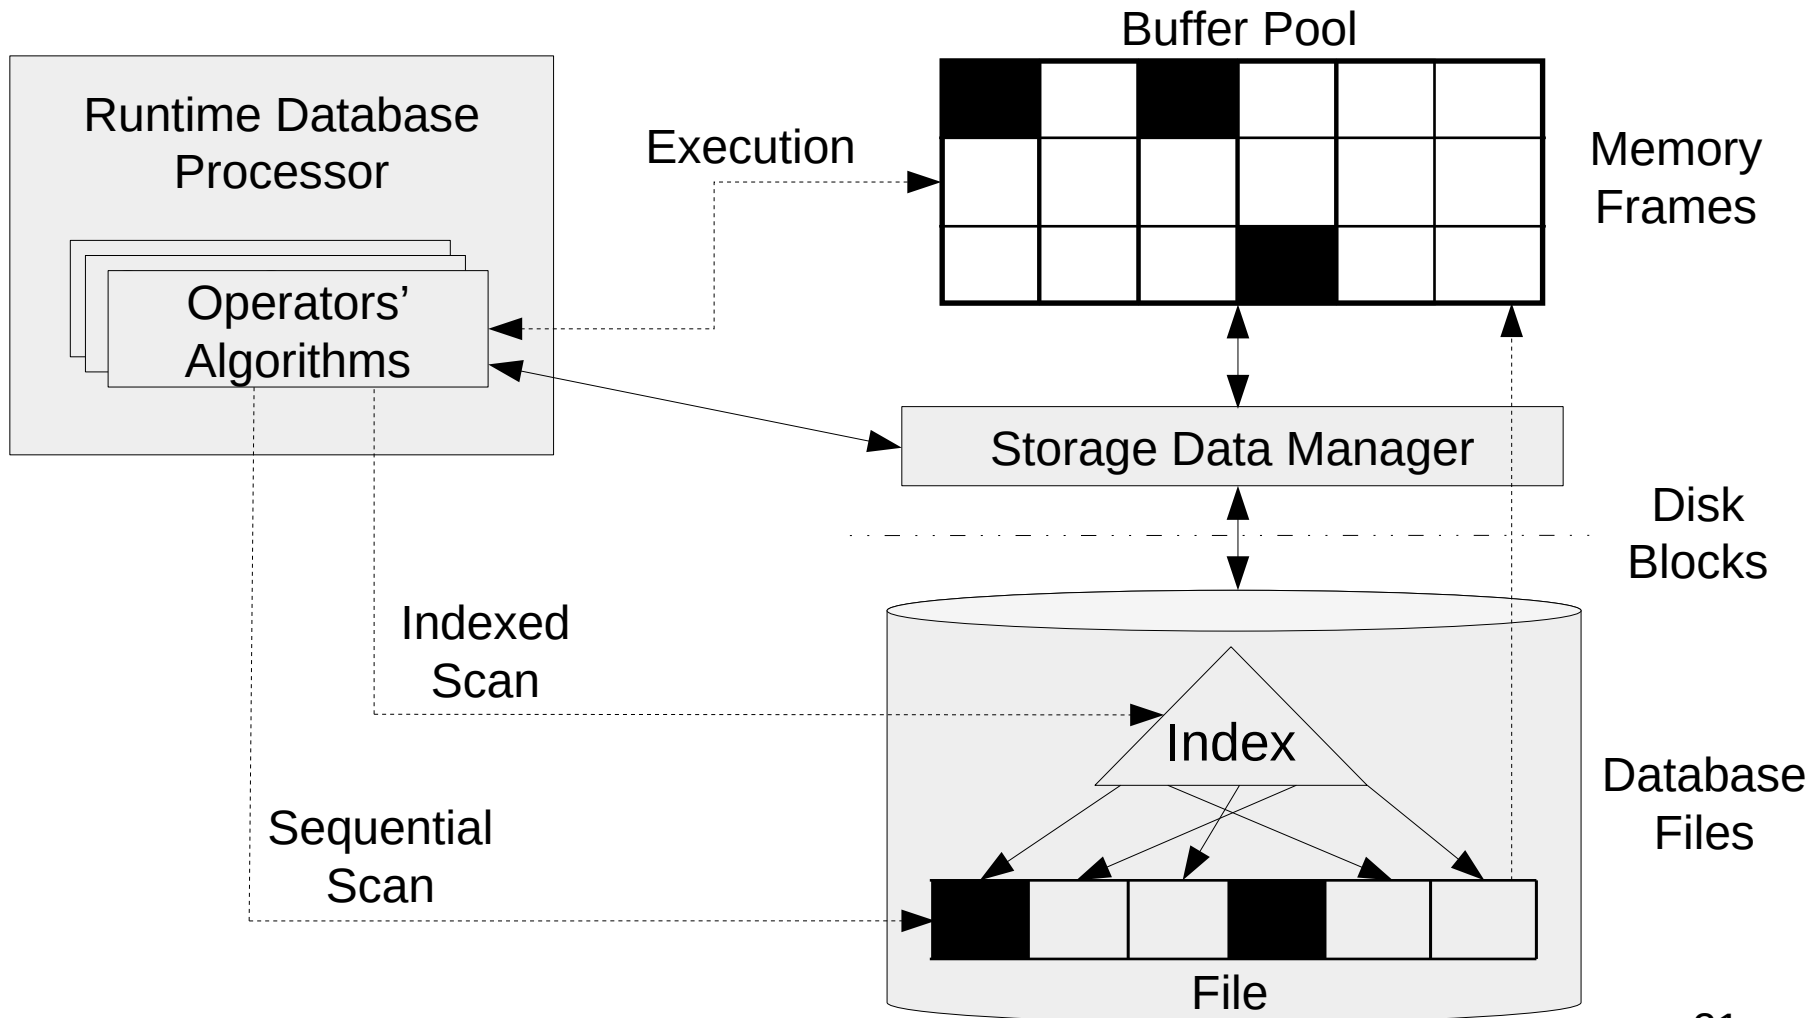
# Paging of Files

- We cannot access a single byte on a disk
  - To access a byte, we have to access a block
- Files are organized in blocks/pages to benefit from this
  - The database files and the data structures that handle them are designed to reduce block accesses

# Placing Records in Files

- Record: collection of related data values or items
  - Values correspond to record field
  - Example: a row in a table
- Large objects
  - Unstructured objects (BLOBs) or semistructured objects (CLOBs)
  - Examples: an image file or a JSON file
- Allocated to disk blocks in both cases
  - Blocking factor
    - Average number of records per block for the file

# The Storage Data Manager



# Buffer Pool

- Memory buffers are also organized in fixed-size blocks
- DBMSs use different memory buffers to store different data
  - Data block buffers, log buffers, etc.
- Different buffers may use distinct block sizes and/or block substitution policies

# Buffer Management and Replacement Strategies

- Buffer management information
  - Pin count
    - To pin a frame in the buffer pool
  - Dirty bit
    - To indicate a modified block
- Buffer replacement strategies
  - To free frames in the buffer pool
  - Strategies
    - Least recently used (LRU)
    - First-in-first-out (FIFO)

# Select – Sequential Scan

- Operator algorithms usually rely on paged accesses

```
seqScanSelect(Relation R, Condition condition) {  
    for each block b in R do  
        for each tuple r in b do {  
            if (condition(r)){  
                result.add(r);  
            }  
        }  
    }  
    return result;  
}
```

- The Storage Data Manager apply techniques to speed up the proccess, including *caching*, *prefetching*, *bulk access*, and others



# Select – Indexed Scan

- The indexed scan is employed when there is an index on one or more attributes in the select condition
  - There are different algorithms according to the index types available

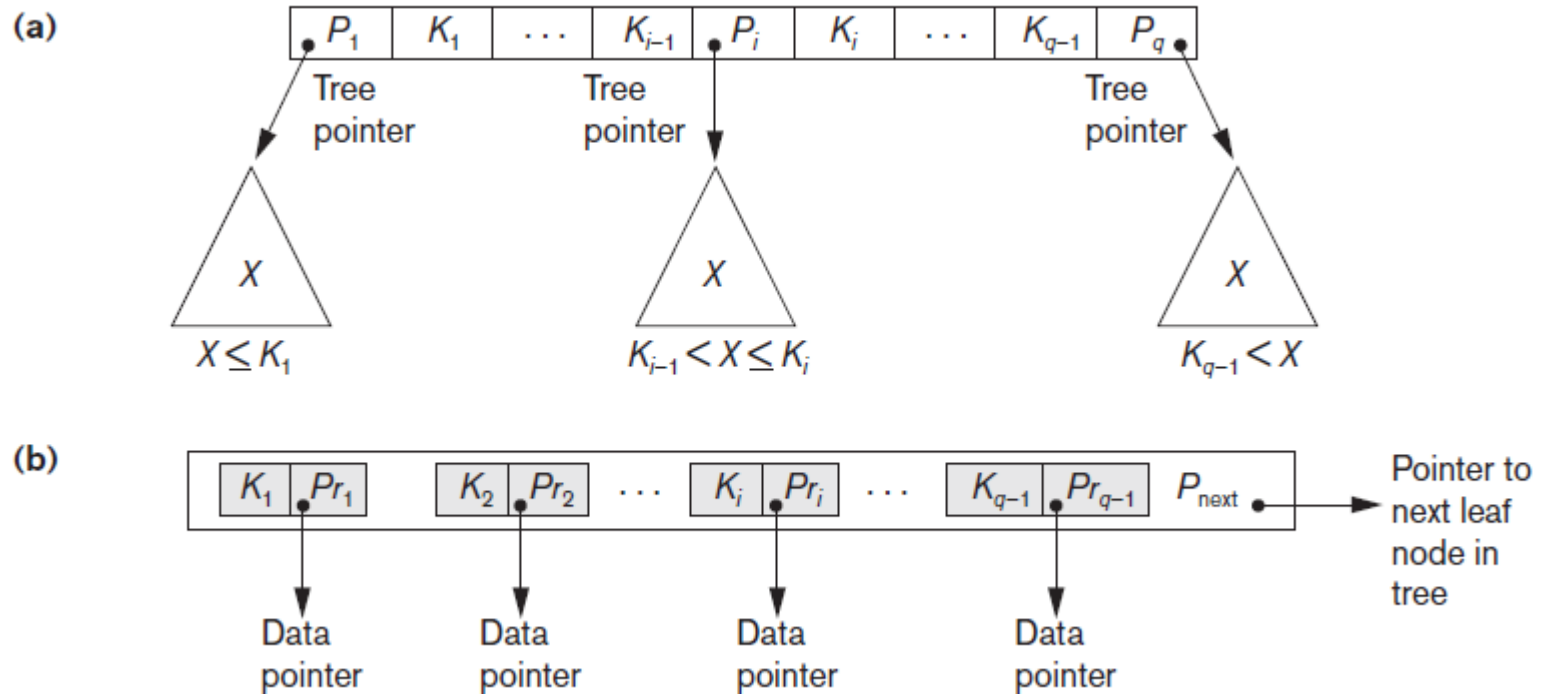
# Secondary and Clustering Indexes

- Secondary index
  - The most common type of indexes
  - A table can have several secondary indexes
    - Syntax: `CREATE [UNIQUE] INDEX my_ix ON table(attribute [<, attribute2, ...>]);`
- Clustering index
  - Sorts physically the data file according to the indexing attribute and creates the index
    - To benefit from the block factor in range queries
    - The drawback is the overhead to maintain the file sorted
  - One clustering index per table
    - Syntax: `CREATE [UNIQUE] CLUSTERING INDEX my_ix ON table(attribute [<, attribute2, ...>]);`


# The B+ -Tree Dynamic Multilevel Index

- Disk-based search tree
  - Nodes have the size of a block
  - Given two search keys  $K_{i-1} < K_i$ , the elements stored in the corresponding subtree are  $K_{i-1} < X \leq K_i$
  - Tree pointers are pointers to blocks: data file + block
  - Record pointers (rids) have the structure: data file + block + record
- Balanced tree
  - Reorganized at each insert or delete using split and unsplit of nodes
  - Nodes are also at least half full, except the root node
  - Bottom-up construction

# B+ -Tree: Node Types




# Secondary Indexes




# Indexed Scan using Secondary Indexes

- Unique scan: WHERE attribute = value (attribute is key)

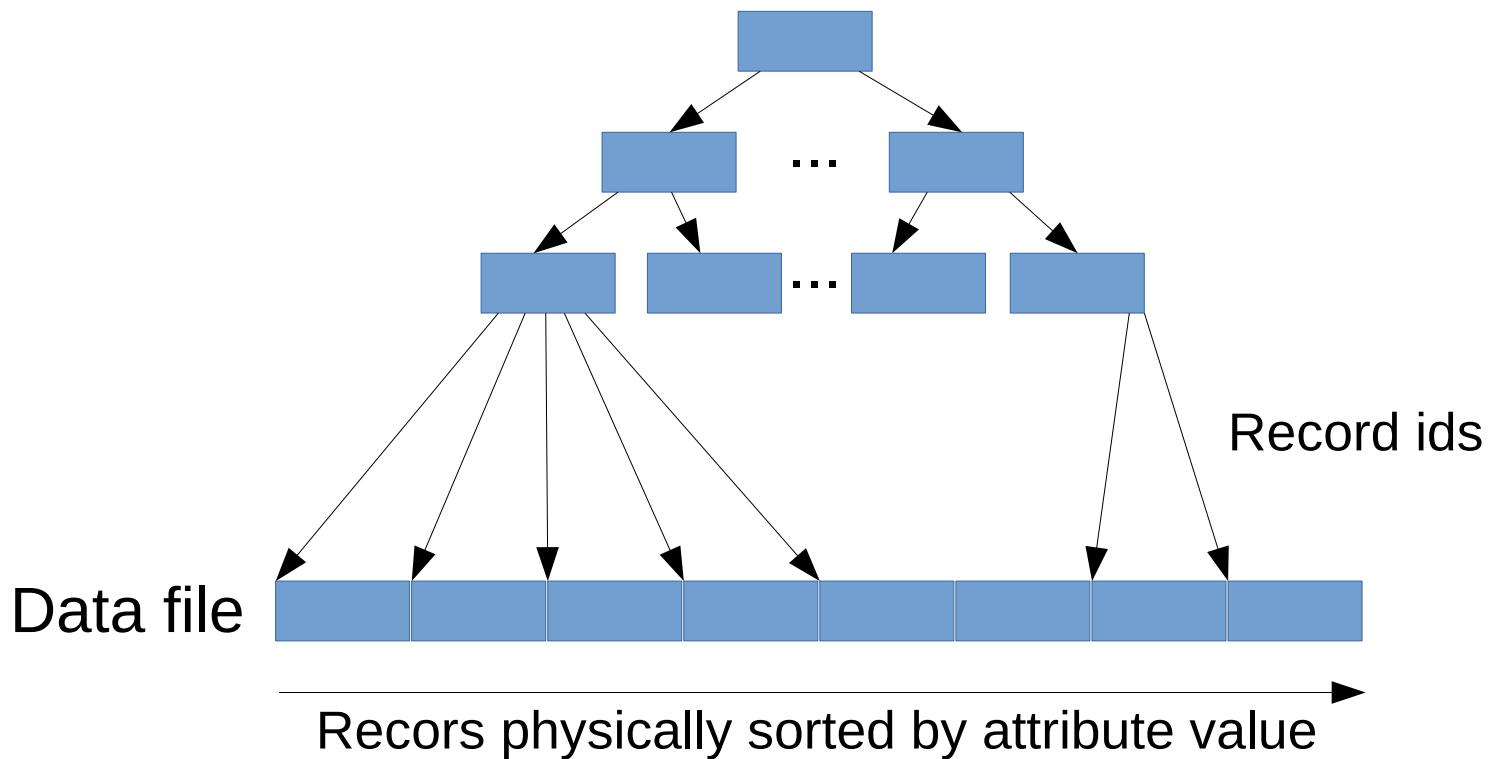


# Indexed Scan using Secondary Indexes (2)

- Range scan: WHERE attribute = value (attribute is not key)  
WHERE attribute >= value, etc.



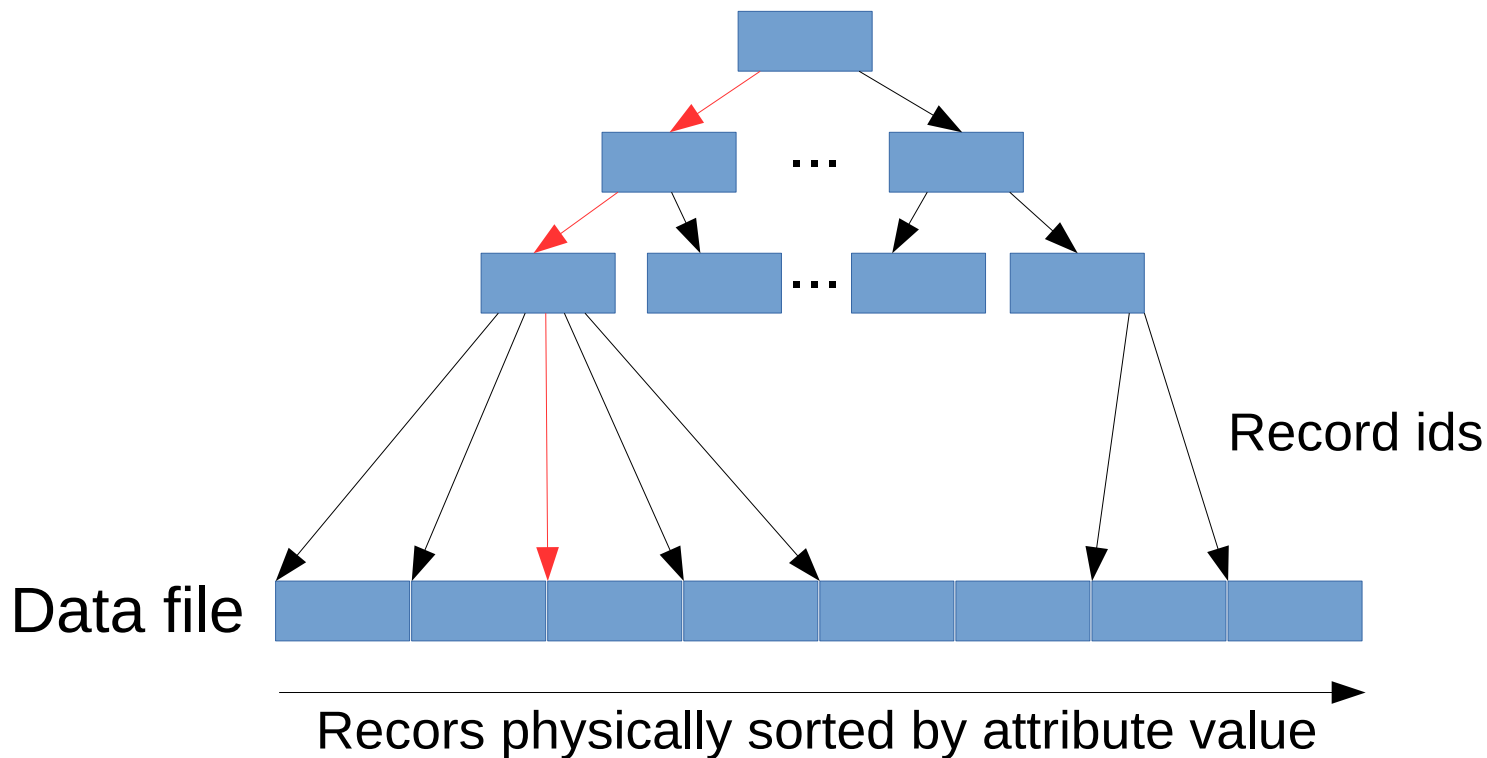
# Clustered Indexes





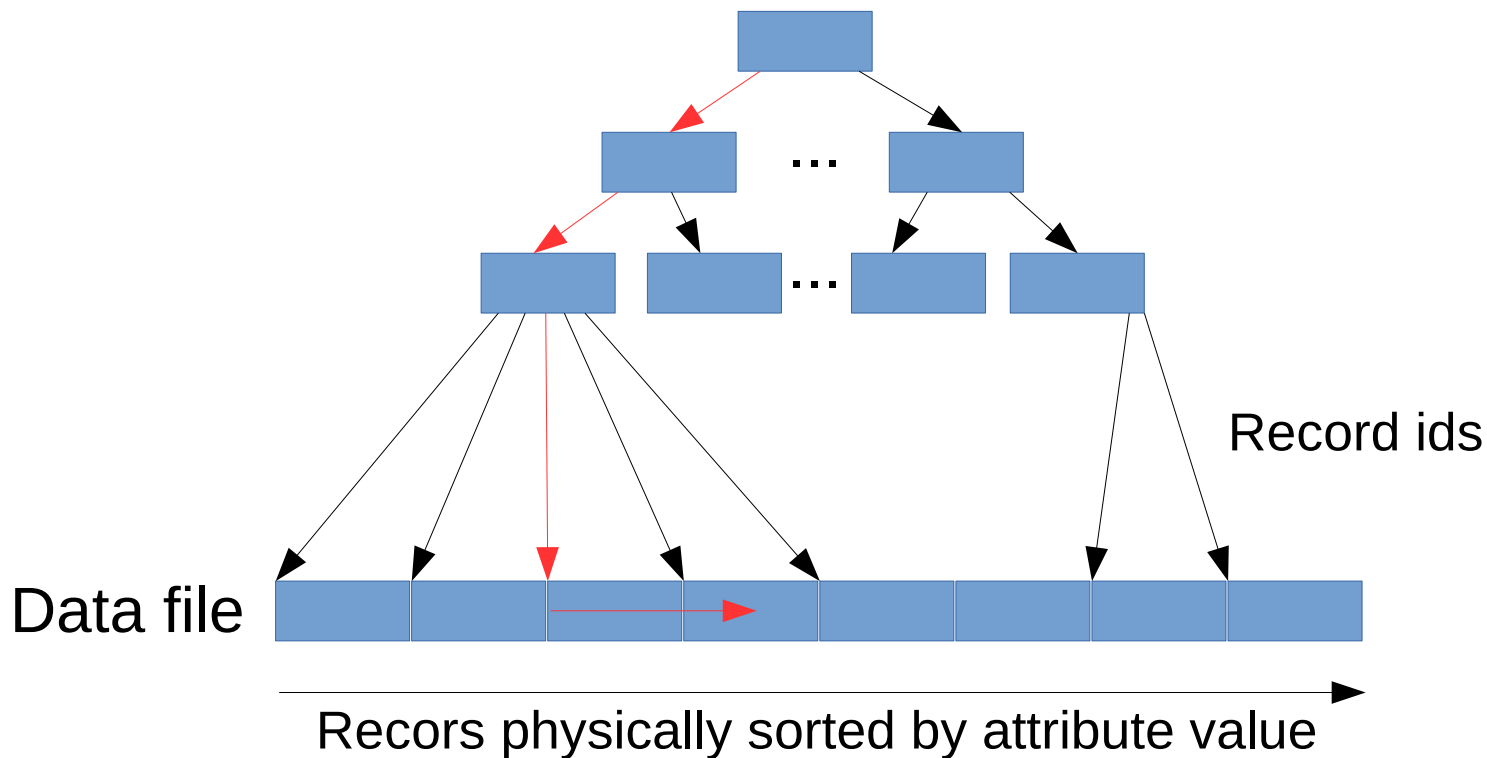
# Indexed Scan using Clustered Indexes

- Unique scan: WHERE attribute = value (attribute is key)



# Indexed Scan using Clustered Indexes (2)

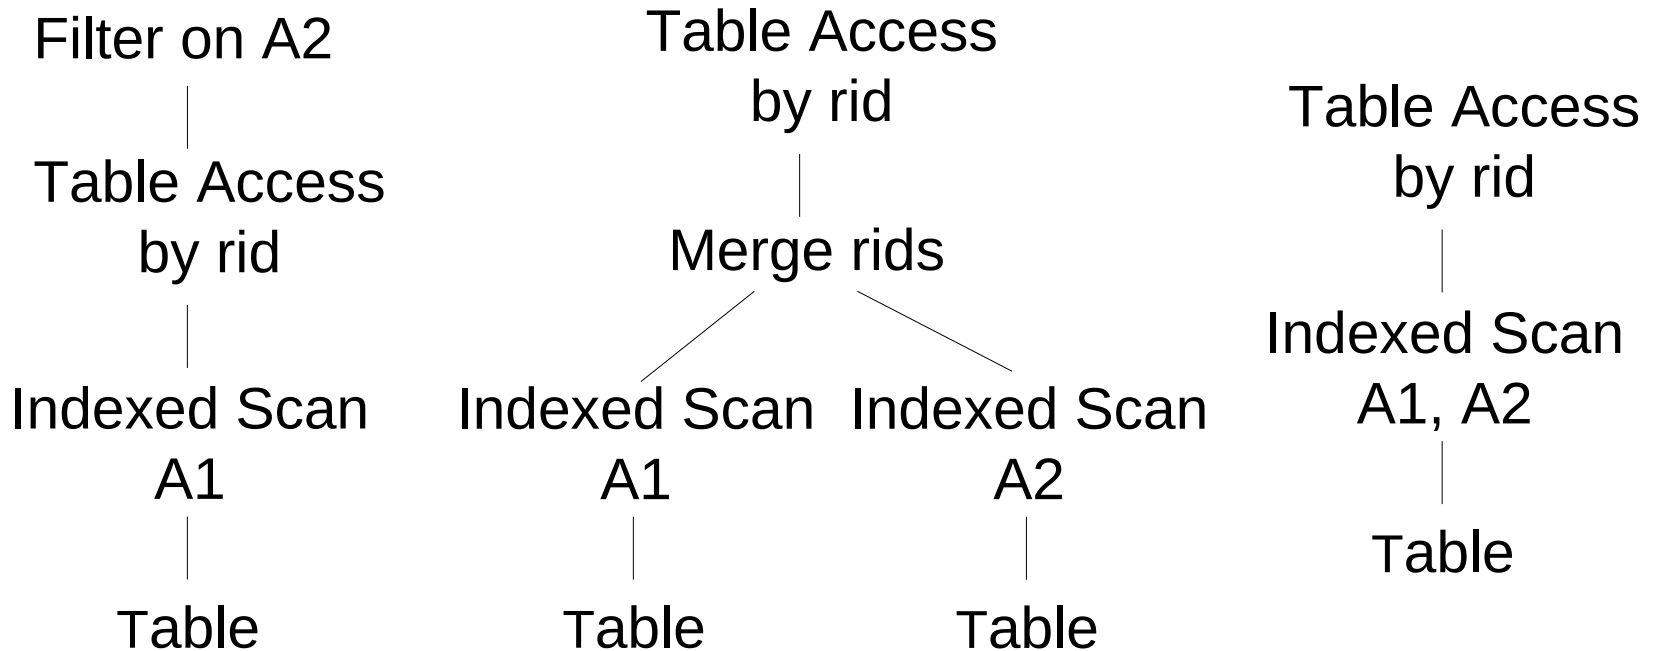
- Range scan: WHERE attribute = value (attribute is not key)  
WHERE attribute >= value, etc.



# Select with Composite Conditions

- The select operator's condition may involve more than one attribute from the same table

WHERE  $t.A1 = v1$  AND  $t.A2 = v2$




# How to Decide the Algorithm to Choose?

- Information stored in DBMS catalog and used by optimizer to compute the cost of the algorithms
  - Tables: # of rows, # of disk blocks, average record size, organization (e.g., sorted or not sorted), etc.
  - Indices: # of levels, # of leaf blocks, uniqueness, etc.
  - Attributes: # of distinct values, minimum and maximum values, etc.

# Attribute Selectivity

- Allows calculation of selection cardinality
  - Estimated number of records that satisfy a selection condition on that attribute
    - Examples:  $\sigma_{\text{salary} = 120k}(\text{Emp})$  ;  $\sigma_{\text{salary} \geq 120k}(\text{Emp})$
- Statistics:
  - Minimum, maximum and # of distinct values
  - Histogram



# Transaction Processing

- A database is a shared resource accessed by many users and processes concurrently
- Not managing this concurrent access to a shared resource will cause problems
  - Problems due to concurrency
  - Problems due to failures
- Transaction processing comprehends techniques to manage a database in a concurrent multi-user environment

# Transactions

- A transaction (sequence of executing operations) may be:
  - Stand-alone, specified in a high level language like SQL submitted interactively, or
  - More typically, embedded within application program
- Transaction boundaries: `Begin_transaction` and `End_transaction`
  - Application program may include specification of several transactions separated by `Begin` and `End` transaction boundaries
  - Transaction code can be executed several times (in a loop), spawning multiple transactions
- Transactions can end in two states:
  - Commit: transaction successfully completes and its results are committed (made permanent)
  - Abort: transaction does not complete and none of its actions are reflected in the database

# What Can Go Wrong?

## T1 at ATM window #1

- 1 read\_item(savings);
- 2 savings = savings - \$100;
- 3 write\_item(savings);
- 4 read\_item(chequing);
- 5 chequing = chequing + \$100;
- 6 write\_item(chequing);

## T2 at ATM window #2

- a read\_item(chequing);
- b chequing = chequing - \$20;
- c write\_item(chequing);
- d dispense \$20 to customer;

- System might crash after transaction begins and before it ends:
  - Money lost if between 3 and 6 or between c and d
  - Updates lost if write to disk not performed before crash
- Chequing account might have incorrect amount recorded:
  - \$20 withdrawal might be lost if T2 executed between 4 and 6
  - \$100 deposit might be lost if T1 executed between a and c
    - In fact, same problem if just 6 executed between a and c



# ACID Properties

- Atomicity

- A transaction is an atomic unit of processing; it is either performed in its entirety or not performed at all

- Consistency

- A correct execution of the transaction must take the database from one consistent state to another

- Isolation

- Even though transactions are executing concurrently, they should appear to be executed in isolation
- That is, their final effect should be as if each transaction was executed in isolation from start to finish

- Durability

- Once a transaction is committed, its changes (writes) applied to the database must never be lost because of subsequent failure

# ACID Properties (2)


- Enforcement of ACID properties
  - Consistency (and application program correctness)
    - Responsibility of the database constraint system
  - Isolation
    - Responsibility of the concurrency control mechanism
  - Atomicity and Durability
    - Ensured by the recovery system

# Transaction Processing Model

- Simple database model:
  - Database: collection of named data items
- Granularity (size) of each data item immaterial
  - A field (data item value), a record, or a disk block
  - TP concepts are independent of granularity
- Basic operations on an item X:
  - read\_item(X): Reads a database item X into a program variable
    - For simplicity, assume that the program variable is also named X
  - write\_item(X): Writes the value of program variable X into the database item named X
- Read and write operations take some amount of time to execute


# Read and Write Operations

- **read\_item(X)** includes the following steps:
  - Find the address of the disk block that contains item X
  - Copy that disk block into a buffer in main memory (if that disk block is not already in some main memory buffer)
  - Copy item X from the buffer to the program variable named X



# Read and Write Operations (2)

- **write\_item(X)** includes the following steps:
  - Find the address of the disk block that contains item X
  - Copy that disk block into a buffer in main memory (if it is not already in some main memory buffer)
  - Copy item X from the program variable named X into its correct location in the buffer
  - Store the updated block from the buffer back to disk
    - Either immediately or, more typically, at some later point in time



# Concurrency Control Protocols

- Two-phase locking protocols
  - Lock data items to prevent concurrent access
- Timestamp ordering protocols
  - Assign a unique identifier for each transaction
  - Apply rules to control how transactions access items according to the timestamps
- Multiversion techniques
  - Kept several versions of an item
  - Accept some read operations that would be rejected in other techniques by reading an older version of the item while maintaining serializability
- Optimistic techniques
  - Perform no checking is done while the transaction is executing
  - Execute a validation phase to check whether any of transaction's updates violate serializability, and commit or abort transactions based on result

# Failure Recovery

- Several approaches exist each using some subset of recovery concepts
  - Write-ahead logging
  - In-place versus shadow updates
  - Immediate versus deferred update

# The UNDO-REDO Approach

- After a crash:
  - 1) UNDO transactions that had not committed
    - To ensure Atomicity
    - Partial results of uncommitted transactions should be discarded
  - 2) REDO committed transactions
    - To ensure Durability
    - Modified blocks in memory might have not be written back to disk
- Uses a system log with write-ahead logging policy and checkpointing
- Example: the ARIES recovery algorithm, used in many IBM relational database products



# System Log

- Append-only file
  - Keep track of all operations of all transactions
  - In the order in which operations occurred
- Stored on disk
  - Persistent except for disk or catastrophic failure
  - Periodically backed up
  - Guard against disk and catastrophic failures
- Main memory buffer
  - Holds records being appended
  - Occasionally whole buffer appended to end of log on disk (flush)

# System Log Entries

- [start\_transaction, T]
  - Transaction T has started execution.
- [write\_item, T, X, old\_value, new\_value]
  - T has changed the value of item X from old\_value to new\_value.
  - Before Image (old\_value) needed to undo(X)
  - After Image (new\_value) needed to redo(X)
- [commit, T]
  - T has completed successfully and committed
  - T's effects (writes) must be durable
- [abort, T]
  - T has been aborted
  - T's effects (writes) must be ignored and undone


# Write-Ahead Logging (WAL)

- Used to ensure that the log is consistent with the database, and to ensure that the log can be used to recover the database to a consistent state
- Two rules:
  - Log record for a page must be written before corresponding page is flushed to disk
    - For Atomicity, so that each operation is known and can be undone if necessary
  - All log records must be written before commit
    - For Durability, so that the effect of a committed transaction is known
- A transaction is said to be committed (reached the commit point) when:
  - All of its operations are executed, and
  - All its log records are flushed to disk

# Checkpointing

- Employed to save redo effort
- Occasionally flush data buffers
  - 1) Suspend execution of transactions temporarily
  - 2) Force-write modified (dirty) buffer data to disk
  - 3) Append [checkpoint] record to log
  - 4) Flush log to disk
  - 5) Resume normal transaction execution
- During recovery, redo required only for log records appearing after [checkpoint] record

# Transaction Status for Recovery



## ■ UNDO/REDO

- $T_1$ : do nothing
- $T_2, T_3$ : REDO
- $T_4, T_5$ : UNDO

# The UNDO/REDO Recovery Process

## 1) ANALYSIS

- a) Scan log from the checkpoint
- b) Identify the transactions to be redone and undone

## 2) REDO (roll-forward)

- a) Scan the log from head to tail (forwards in time)
- b) Redo updates of committed transactions
- c) Use after image for new values

## 3) UNDO (roll-back)

- a) Scan log from tail to head (backward in time)
- b) Restore before image by undoing updates of active transactions, writing compensating log records for undone operations
  - e.g., the compensation action for an insert is a delete
- c) In case of crash during recovery, the algorithm redoes up to the last compensation operation and starts undoing from the previous one